


# Chapter 10 of AEPV

*DJM, Revised NAK*

*26 February 2019*

## Assumptions on Residuals



```
##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.164973  0.4980974  4.346484 3.383817e-05
## x           2.127688  0.1853945 11.476538 7.852372e-20
```

Questions: Whats wrong here? Where can we most reliably estimate the linear function? Where is the most unreliable area?

## Standard Errors in Heteroskedacity


```
ols.heterosked.example = function(n) {
  y = 3 - 2 * x + rnorm(n, 0, sapply(x, function(x) {
    1 + 0.5 * x^2
  })))
  fit.ols = lm(y ~ x)
  return(fit.ols$coefficients - c(3, -2))
}
```

```
ols.heterosked.error.stats = function(n, m = 10000) {
  ols.errors.raw = t(replicate(m, ols.heterosked.example(n)))
  intercept.se = sd(ols.errors.raw[, "(Intercept)"])
  slope.se = sd(ols.errors.raw[, "x"])
  return(c(intercept.se = intercept.se, slope.se = slope.se))
}
```

```
ols.heterosked.error.stats(100)
```

```
## intercept.se      slope.se
##    0.6781456      0.5248545
```

## Another One



```
##           Estimate Std. Error t value    Pr(>|t|)
## (Intercept) 2.563258  0.7385148  3.470828 0.0007731719
## x           2.494290  0.6341567  3.933239 0.0001565225
```

What about in this graph? Where is more reliable for estimating the line, and where is less reliable?

## Ordinary Least Squares: A review

In Ordinary Least Squares, we are trying to minimize the sum of squared errors:

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n (y_i - \underline{x}_i^\top \beta)^2 = (X^\top X)^{-1} X^\top Y$$

The hat matrix is

$$\hat{Y} = X\hat{\beta} = X(X^\top X)^{-1} X^\top Y = HY$$

The Gauss-Markov theorem says if:

1.  $Y_i = \underline{x}_i^\top \beta + \epsilon_i$
2.  $\mathbb{E}[\epsilon_i] = 0$
3.  $\mathbb{V}[\epsilon_i] = \sigma^2 < \infty$
4.  $\operatorname{Cov}[\epsilon_i, \epsilon_j] = 0$

Then  $\hat{\beta} = (X^\top X)^{-1} X^\top Y$  has the smallest variance of all possible unbiased estimators for  $\beta$ .

In linear models theory, we call the line based off of  $\hat{\beta}$  the **Best Linear Unbiased Estimator (BLUE)** of the true line/coefficients.

## Weighting in the Least Squares formula

Weighted least-squares (WLS) is based on the following:

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n w_i (y_i - \underline{x}_i^\top \beta)^2 = (X^\top W X)^{-1} X^\top W Y$$

- If some of those assumptions for G-M are violated, in particular, if  $\mathbb{V}[\epsilon_i]$  depends on  $x_i$  (notated like  $\sigma^2(x_i)$ ), then we lose the optimality of OLS.
- Aside: Gauss-Markov is a commonly used justification for OLS in applied work. The logic goes like this: (1) unbiased is good, (2) G-M says OLS is the best linear model which is unbiased. The problem is that (1) is wrong. Unbiased may be good, but often a little bias is better.

The main question here is, how do we choose the weights?  $w_i = ??$

## Choosing $w_i$

Lets consider the data from the first slide.


We could set  $w_i = 0$  for a certain range of observations, and  $w_i$  for other observations.

For example, use  $w_i = 1$  for  $|x_i| \leq 2$  and  $w_i = 0$  otherwise.

$$\hat{\beta} = \begin{pmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \end{pmatrix} = \underset{\beta_0, \beta_1}{\operatorname{argmin}} \sum_{i \in S} (y_i - (\beta_0 + \beta_1 x_i))^2, \text{ where } S = \{i : |x_i| \leq 2\}$$

```
#Make weights
w = rep(NA,n)
w[abs(dfHetero$x) <= 2] <- 1
w[abs(dfHetero$x) > 2] <- 0
dfHetero$w <- w
wlm1 <- lm(y ~ x, dfHetero, weights=w) # For Next Part
```

```
ggplot(dfHetero, aes(x,y)) + geom_point() +
  geom_smooth(method='lm', se = F, color = red) +
  geom_smooth(method='lm', aes(x = x, y = y, weight=w), se=FALSE,color=blue)
```



Is that any better?

## Checking The Models (1)

Model 1 Summary:

```
##
## Call:
## lm(formula = y ~ x, data = dfHetero)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-18.8477	-1.3648	0.5328	1.9607	16.3377

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	2.1650	0.4981	4.346	3.38e-05
x	2.1277	0.1854	11.477	< 2e-16

```
##
## Residual standard error: 4.971 on 98 degrees of freedom
## Multiple R-squared: 0.5734, Adjusted R-squared: 0.569
## F-statistic: 131.7 on 1 and 98 DF, p-value: < 2.2e-16
```

### Weighted Model 1 Summary:

```
##
## Call:
## lm(formula = y ~ x, data = dfHetero, weights = w)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -4.299 -0.479  0.000  0.000  3.498
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.1801     0.2251   14.13 < 2e-16
## x             2.2183     0.1962   11.31 5.7e-16
##
## Residual standard error: 1.691 on 55 degrees of freedom
## Multiple R-squared:  0.6992, Adjusted R-squared:  0.6937
## F-statistic: 127.9 on 1 and 55 DF,  p-value: 5.704e-16
```

### Choosing $w_i$ (Part 2)


Now lets look at the other data from the beginning

Use  $w_i = 1$  for  $x_i < 1.2$  and  $w_i = 0$  otherwise.

$$\hat{\beta} = \begin{pmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \end{pmatrix} = \underset{\beta_0, \beta_1}{\operatorname{argmin}} \sum_{i \in S} (y_i - (\beta_0 + \beta_1 x_i))^2, \text{ where } S = \{i : x_i < 1.2\}$$

```
#Make weights
w = rep(NA,n)
w[dfHetero2$x < 1.2] <- 1
w[dfHetero2$x >= 1.2] <- 0
dfHetero2$w <- w
wlm2 <- lm(y ~ x, dfHetero2, weights=w) # For Next Part

ggplot(dfHetero2, aes(x,y)) + geom_point() +
  geom_smooth(method='lm', se = F, color = red) +
  geom_smooth(method='lm', aes(x = x, y = y, weight=w), se=FALSE,color=blue)
```



## Checking The Models (Part 2)

### Model 2 Summary:

```
##
## Call:
## lm(formula = y ~ x, data = dfHetero2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -7.8547 -1.6219 -0.1132  1.5305 11.4908
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.5633     0.7385   3.471 0.000773
## x             2.4943     0.6342   3.933 0.000157
##
## Residual standard error: 3.377 on 98 degrees of freedom
## Multiple R-squared:  0.1363, Adjusted R-squared:  0.1275
## F-statistic: 15.47 on 1 and 98 DF,  p-value: 0.0001565
```

### Weighted Model 2 Summary:

```
##
## Call:
## lm(formula = y ~ x, data = dfHetero2, weights = w)
##
```

```
## Weighted Residuals:
##      Min      1Q  Median      3Q      Max
## -3.0455 -0.3595  0.0000  0.0081  4.6723
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.1810     0.5418   4.026 0.000175
## x              3.4659     0.7677   4.515 3.39e-05
##
## Residual standard error: 1.779 on 55 degrees of freedom
## Multiple R-squared:  0.2704, Adjusted R-squared:  0.2571
## F-statistic: 20.38 on 1 and 55 DF,  p-value: 3.392e-05
```

## The General Issue With Heteroskedacity

So suppose  $\mathbb{V}[\epsilon_i] = \sigma^2(x_i)$ . That is our “homoskedasticity” assumption is violated. Should we care?

What if we just use OLS (that is `lm`) anyway?

Some things don’t change.

1. We still have that  $\mathbb{E}[\hat{\beta}] = \beta$ . That is OLS is still unbiased.
2. We still have that OLS minimizes the sum of squared residuals: among all lines, OLS makes  $\sum_{i=1}^n (x_i^\top \hat{\beta} - y_i)^2$  as small as possible.

Some things **do** change.

1. OLS no longer has the best variance of all unbiased estimators (WLS does).
2. The standard errors that R produces are wrong. They make it seem “more certain” than is correct (could use the bootstrap to fix it though).
3. So are the  $F$ -tests and  $p$ -values (again, the bootstrap).

## Optimal WLS for Heteroskedacity: $\sigma^2(x)$

So WLS is fairly general. But for now, let’s focus on how to use it for heteroskedasticity.

Suppose you **know** the following:

1.  $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$ .
2.  $\mathbb{E}[\epsilon_i] = 0$
3.  $\mathbb{V}[\epsilon_i] = \sigma^2(x_i)$  ( $\sigma^2(\cdot)$  is a function).

It can be shown that the optimal weights are  $w_i = \frac{1}{\sigma^2(\underline{x}_i)}$ , making no assumption about the probability distribution of the errors, besides what is above.

This means, that the optimal  $\hat{\beta}$  vector is found by minimizing

$$\sum_{i=1}^n \frac{(y_i - \underline{x}_i^\top \hat{\beta})^2}{\sigma_i^2(\underline{x}_i)}$$

See section 10.2.2.1 of Shalizi’s book if you are curious... (Actually, I don’t recommend that.)

## Weighting in Kernel Regression, an aside

Try to recall linear smoothers, and Kernel Regression in particular.

Kernel Regression can be written as a sort of Weighted Least Squares solution

$$\hat{c} = \underset{c}{\operatorname{argmin}} \sum_{j=1}^n \sum_{i=1}^n w_{ij} (y_i - c_j)^2 \quad w_{ij} = \frac{K((x_i - x_j)/h)}{\sum_{i=1}^n K((x_i - x_j)/h)}$$

This is locally constant regression.


You don't need to understand this formula, but it can be useful, and it provides some justification for WLS based on previous ideas.

## Using Optimal Weights in LM

```
#Make weights

dfHetero$w <- 1/(1+x^2/2) # First example had Var(residuals) = (1+x^2/2)
opt.wlm1 <- lm(y ~ x, dfHetero, weights=w) # For Next Part

with(dfHetero, plot(x,y))
abline(opt.wlm1, col = red)
abline(lm1, col = blue)
```



## Comparing the different models

Model 1 Summary:

```
##
## Call:
```



```
## lm(formula = y ~ x, data = dfHetero)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -18.8477  -1.3648   0.5328   1.9607  16.3377
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.1650     0.4981   4.346 3.38e-05
## x             2.1277     0.1854  11.477 < 2e-16
##
## Residual standard error: 4.971 on 98 degrees of freedom
## Multiple R-squared:  0.5734, Adjusted R-squared:  0.569
## F-statistic: 131.7 on 1 and 98 DF,  p-value: < 2.2e-16
```

---

\_\_\_ Pseudo Weights Model 1 Summary\_\_\_:

```
##
## Call:
## lm(formula = y ~ x, data = dfHetero, weights = w)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -4.299 -0.479   0.000   0.000   3.498
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   3.1801     0.2251  14.13 < 2e-16
## x             2.2183     0.1962  11.31 5.7e-16
##
## Residual standard error: 1.691 on 55 degrees of freedom
## Multiple R-squared:  0.6992, Adjusted R-squared:  0.6937
## F-statistic: 127.9 on 1 and 55 DF,  p-value: 5.704e-16
```

---

Optimal Weights Model 1 Summary:

```
##
## Call:
## lm(formula = y ~ x, data = dfHetero, weights = w)
##
## Weighted Residuals:
##      Min       1Q   Median       3Q      Max
## -16.8430  -0.9417   0.6378   1.8065  13.9880
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.9154     0.5136   3.729 0.000322
## x             2.1391     0.1926  11.105 < 2e-16
##
## Residual standard error: 4.16 on 98 degrees of freedom
## Multiple R-squared:  0.5572, Adjusted R-squared:  0.5527
## F-statistic: 123.3 on 1 and 98 DF,  p-value: < 2.2e-16
```

## Simulating the difference between OLS and WLS

We will compare the simulated OLS and WLS Standard Errors from the model at the very beginning

```
## intercept.se      slope.se
##      0.3066330     0.3553108

## intercept.se      slope.se
##      0.2734971     0.3190810
```

If we knew  $\sigma^2(x)$ , this would be easy as... pie. Unfortunately, it is never that easy.

This means that our new issue is estimating  $\sigma^2(x)$ .

## Variances and Conditional Variances

In general, for a random variable  $X$ , the variance is defined as:

$$\mathbb{V}[X] = \mathbb{E} \left[ (X - \mathbb{E}[X])^2 \right]$$

Let's consider the variance of the residuals:  $\epsilon_i = y_i - \mathbf{x}_i^\top \beta$


$$\begin{aligned} \sigma^2(x_i) &= \mathbb{V}[\epsilon_i | x_i] = \mathbb{E} \left[ \epsilon_i - \mathbb{E}[\epsilon_i | x_i]^2 | x_i \right] \\ &= \mathbb{E} [\epsilon_i^2 | x_i] \end{aligned}$$

What is our estimate of this expectation?

## Estimating $\sigma^2(x)$ , An Iterative Process


1. Use `lm` to estimate  $\beta_0$  and  $\beta_1$  to get the estimated regression line  $\hat{\mu}(x)$ .
2. Use your estimated regression line to calculate the squared residuals,  $e_i^2 = (y_i - \hat{\mu}(x_i))^2$ .
3. Use nonparametric regression to get  $\hat{\sigma}^2(x)$ , which is an estimate of  $\mathbb{E}[\epsilon_i^2 | x_i]$
4. Use this estimate “know”  $\sigma^2(x)$  and use WLS (with `lm(y~x, weights=1/sig2)`)
5. You could stop here. But since you now have “better” estimates of  $\beta_1$  and  $\beta_0$ , it's better to iterate 2 and 3 until some convergence.
6. Ok. Something converged, so you return the last estimates of  $\beta_0$  and  $\beta_1$ . But the SEs are not right quite right since we are only estimating  $\sigma^2(x)$
7. To get “correct” SEs, use the bootstrap:
  - a. Non-parametric: repeat 1-5  $B$  times on resampled data. This can be rather slow...
  - b. Model-based: this is actually pretty hard here, better not to do it.

## An Example Using the First Model



```
##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.164973  0.4980974  4.346484 3.383817e-05
## x           2.127688  0.1853945 11.476538 7.852372e-20
```

## Using Kernel Regression to Estimate $\sigma^2(x)$



### Iterations for first model

```
##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.164973  0.4980974  4.346484 3.383817e-05
## x           2.127688  0.1853945 11.476538 7.852372e-20

var1 <- npreg(residuals(lm1)^2 ~ x, data = dfHetero )
wlm1 <- lm(y ~ x, dfHetero, weights = 1/fitted(var1))
summary(wlm1)$coefficients

var2 <- npreg(residuals(wlm1)^2 ~ x, data = dfHetero )
wlm2 <- lm(y ~ x, dfHetero, weights = 1/fitted(var2))
summary(wlm2)$coefficients

var3 <- npreg(residuals(wlm2)^2 ~ x, data = dfHetero )
wlm3 <- lm(y ~ x, dfHetero, weights = 1/fitted(var3))
summary(wlm3)$coefficients

##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.878506  0.2861681 10.05879 9.021827e-17
## x           2.013979  0.1738771 11.58277 4.651842e-20

##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.907036  0.2757588 10.54195 8.093778e-18
## x           2.009604  0.1733555 11.59239 4.436706e-20

##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.907659  0.2755368 10.55271 7.671609e-18
## x           2.009335  0.1733433 11.59165 4.452920e-20
```


## Simplified Iterative Function

```
iterative.wls <- function(x, y, tol = 0.01, max.iter = 100) {  
  iteration <- 1  
  old.coefs <- NA  
  regression <- lm(y ~ x)  
  coefs <- coefficients(regression)  
  while (is.na(old.coefs) || ((max(coefs - old.coefs) > tol) && (iteration <  
    max.iter))) {  
    variance <- npreg(residuals(regression)^2 ~ x)  
    old.coefs <- coefs  
    iteration <- iteration + 1  
    regression <- lm(y ~ x, weights = 1/fitted(variance))  
    coefs <- coefficients(regression)  
  }  
  return(list(regression = regression, variance = variance, iterations = iteration))  
}
```

## Alternative, log of residuals

1. Use `lm` to estimate  $\beta_0$  and  $\beta_1$ .
2. Use your estimated regression line to calculate the squared residuals,  $e_i^2 = (y_i - \hat{\mu}(x_i))^2$ .
3. Calculate  $\log(\hat{e}_i^2)$  and use `npreg` to estimate  $\log \sigma^2(x)$ .
4. Now pretend that you “know”  $\sigma^2(x)$  (take exp of your estimate from 2.) and use WLS (with `lm(y~x, weights=1/sig2)`)
5. You could stop here. But since you now have “better” estimates of  $\beta_1$  and  $\beta_0$ , it’s better to iterate 2 and 3 until some convergence.
6. Ok. Something converged, so you return the last estimates of  $\beta_0$  and  $\beta_1$ . But the SEs are not right (because you “know”  $\sigma^2(x)$  but you don’t **know** it).
7. To get SEs, use the bootstrap:
  - a. Non-parametric: repeat 1-5  $B$  times on resampled data. Still slow.
  - b. Model-based: this is actually pretty hard here, better not to do it.

Looking at Log of  $e_i^2$



Same Example, Now using the log

```
logvar1 <- npreg(log(residuals(lm1)^2) ~ x, data = dfHetero )
wlm1 <- lm(y ~ x, dfHetero, weights = 1/exp(fitted(logvar1)))
summary(wlm1)$coefficients

logvar2 <- npreg(log(residuals(wlm1)^2) ~ x, data = dfHetero )
wlm2 <- lm(y ~ x, dfHetero, weights = 1/exp(fitted(logvar2)))
summary(wlm2)$coefficients

logvar3 <- npreg(log(residuals(wlm2)^2) ~ x, data = dfHetero )
wlm3 <- lm(y ~ x, dfHetero, weights = 1/exp(fitted(logvar3)))
summary(wlm3)$coefficients
```

```
##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 2.972818  0.2454297 12.11271 3.458792e-21
## x           2.059395  0.1665938 12.36177 1.027915e-21

##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 3.028201  0.2277862 13.29405 1.153945e-23
## x           2.033225  0.1690139 12.02993 5.183114e-21

##           Estimate Std. Error  t value    Pr(>|t|)
## (Intercept) 3.027455  0.2278962 13.28436 1.208497e-23
## x           2.029709  0.1684034 12.05266 4.637998e-21
```

## A Bigger Example

This is a (slightly modified) portion of a real job interview.

It is a very simple application of heteroskedasticity.

Heteroskedasticity appears frequently with financial data, so those companies like to see if you can handle it.

## The set up

The dataset `jobInt` contains data from a simple linear model with heteroskedastic noise.

```
set.seed(02-26-2019)
n=250
x = rnorm(n, sd=1.5)
sigma.x <- function(x) (5*(sin(x)^2)+2)*(x>=0) + (x^2+1)*(x<0)
y = -1+2*x + sigma.x(x)*rnorm(n)
jobInt = data.frame(x=x, y=y)
```

In other words, for  $i = 1, \dots, 250$ ,

$$y_i = \beta_0 + \beta_1 x_i + \sigma(x_i) \epsilon_i \quad \epsilon_i \sim N(0, 1).$$


You know nothing about (the function)  $\sigma(\cdot)$ .

Your goal is to estimate  $(\beta_0, \beta_1)$  as well as possible, and provide a CI.

## How do I do this?


First things first, EDA.

```
ggplot(jobInt, aes(x,y)) + geom_point(color=blue) +
  geom_smooth(method='lm', se=FALSE, color=red)
```




```
basicMod = lm(y~x)
```

## QQ-Plot of Residuals



## Examine Residuals vs. Fitted values



## Code for Iterative WLS using Log of Squared Residuals

This code takes in data and does steps 1-5. It is **not** optimized for speed, but for readability, so run with care.

```
heteroWLS <- function(dataFrame, tol = 1e-4, maxit = 100, track=FALSE){  
  # inputs: a data object, optional: tolerance, max.iterations, and progress tracker (prints)  
  # outputs: estimated betas and weights  
  require(np)  
  ols = lm(y~x, data=dataFrame)  
  b = coefficients(ols)  
  conv = FALSE  
  for(iter in 1:maxit){ # don't let this run forever
```




```

if(conv) break # if the b's stop moving, get out of the loop
logSqResids = log(residuals(ols)^2)
winv = exp(predict(npreg(logSqResids~x, data=dataFrame, tol=1e-2, ftol=1e-2)))
winv[winv < tol] = tol # zero inverse weights are bad, make them small
ols = lm(y~x, weights = 1/winv, data=dataFrame) #weights are 1 / estim.variance
newb = coefficients(ols)
conv.crit = sum((b-newb)^2) # calculate how much b moved
if(track) cat('\n', iter, '/', maxit, ' conv.crit = ', conv.crit) # print progress
conv = (conv.crit < tol) # check if the b's changed much
b = newb # update the coefficient estimates
}
return(list(betas=b, weights = winv, log2resids = log(residuals(ols)^2)))
}

```

## Log of Squared Residuals in OLS Model



## Running Code (Slow...)

```

start.time <-proc.time()[[3]]

resampWLS <- function(dataFrame,...){ # ... means options passed on
  rowSamp = sample(1:nrow(dataFrame), size=nrow(dataFrame), replace=TRUE)
  return(heteroWLS(dataFrame[rowSamp,],...)$betas) # passed things on if desired
}

B = 100 #
alp = .05
origBetas = heteroWLS(jobInt)
time.1 <- proc.time()[[3]] - start.time

bootBetas <- replicate(B, resampWLS(jobInt, maxit=20))

```

```
qq = apply(bootBetas, 1, quantile, probs=c(1-alp/2, alp/2))
CI = cbind(origBetas$betas, 2*origBetas$betas - t(qq))
colnames(CI) = c('coef', rev(colnames(CI)[2:3]))

time.2 <- proc.time()[[3]] - time.1
```

```
# Time to get WLS function to converge on weights
time.1
```

```
## [1] 15.541
```

```
# Time to get bootstrapped CIs
time.2
```

```
## [1] 81.033
```

```
CI
```

```
##           coef      2.5%      97.5%
## (Intercept) -1.002996 -1.521497 -0.7480195
## x           1.959554  1.374641  2.4224998
```

## Log of Squared Residuals in WLS Model

